Retrieval-Augmented Generation RAG Overview
RAG combines document retrieval with large language models.
Step 1: Text from PDF documents is extracted and split into chunks.
Step 2: Chunks are converted into vector embeddings.
Step 3: Embeddings are stored in ChromaDB vector database.
Step 4: User asks a query, retriever fetches top 5 chunks.
Step 5: Google Gemini LLM answers using the retrieved context.